

New Results in Bounded-Suboptimal Search

Maximilian Fickert¹ and Tianyi Gu² and Wheeler Ruml²



Planning as Heuristic Graph Search

Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

heuristic search: a planning approach

Planning as Heuristic Graph Search

Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

heuristic search: a planning approach

planning is a model-based AI method, it models the environment as a state space and finds a sequence of actions that accomplishes some objective

Planning as Heuristic Graph Search

Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

heuristic search: a planning approach

planning is a model-based AI method, it models the environment as a state space and finds a sequence of actions that accomplishes some objective

heuristic search:


$\{\text{states, actions}\} \rightarrow \{V, E\}$

planning problem \rightarrow find a path from s_{init} to $\{s_{goal}\}$

guide graph search by a heuristic estimate of cost-to-goal

Planning as Heuristic Graph Search

heuristic search associates costs with states,
used to guide search



Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

Planning as Heuristic Graph Search

Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

A*: expands the node with minimal f value
returns optimal path
optimal search can take too long!
because it must expand every node with $f < C^*$ ¹

¹How Good is Almost Perfect, Malte Helmert and Gabriele Roger, AAAI, 2008.

Planning as Heuristic Graph Search

Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

A*: expands the node with minimal f value
returns optimal path
optimal search can take too long!
because it must expand every node with $f < C^*$ ¹

What if we don't have time?

¹How Good is Almost Perfect, Malte Helmert and Gabriele Roger, AAAI, 2008.

Alternatives to Optimal Search: Problem Settings

[Introduction](#)

■ [Heuristic Search](#)

■ **[Problem Settings](#)**

■ [Overview](#)

[Bounded Suboptimal](#)

[New Algorithms](#)

[Results](#)

[Conclusions](#)

optimal: minimize solution cost
expand every node with $f < C^*$

greedy: minimize solving time

anytime: incrementally converge to optimal

bounded-suboptimal: minimize time subject to relative cost
bound (factor of optimal)

bounded-cost: minimize time subject to absolute cost bound

contract: minimize cost subject to absolute time bound

utility function: minimize utility function of cost and time

Alternatives to Optimal Search: Problem Settings

Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

optimal: minimize solution cost
expand every node with $f < C^*$

greedy: minimize solving time

anytime: incrementally converge to optimal

bounded-suboptimal: minimize time subject to relative cost
bound (factor of optimal)

bounded-cost: minimize time subject to absolute cost bound

contract: minimize cost subject to absolute time bound

utility function: minimize utility function of cost and time

Introduction

■ Heuristic Search

■ Problem Settings

■ Overview

Bounded Suboptimal

New Algorithms

Results

Conclusions

- Introduction
- Bounded-Suboptimal Search
- New Algorithms
 - DXES
 - RoundRobin
- Results
- Conclusions

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES

■ DPS

■ XES

New Algorithms

Results

Conclusions

Bounded-Suboptimal Search

Bounded-Suboptimal Search: The Problem Setting

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES


■ DPS

■ XES

New Algorithms

Results

Conclusions



Bounded-Suboptimal Search: The Problem Setting

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES


■ DPS

■ XES

New Algorithms

Results

Conclusions



Objective: Find a plan with cost at most ωC^ as fast as possible.*

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES

■ DPS

■ XES

New Algorithms

Results

Conclusions

Three source of heuristic information:

h : a lower bound on cost-to-go

$$f(n) = g(n) + h(n)$$

traditional A* lower bound

\hat{h} : an estimate of cost-to-go

$$\hat{f} = g(n) + \hat{h}(n)$$

unbiased estimates can be more informed

\hat{d} : an estimate of distance-to-go

nearest goal is the easiest to find

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES

■ DPS

■ XES

New Algorithms

Results

Conclusions

Three source of heuristic information: h, \hat{h}, \hat{d}

EES search strategy:

$best_f$: open node giving lower bound on cost

$best_{\hat{f}}$: open node giving estimated optimal cost

$best_{\hat{d}}$: estimated ω -suboptimal node with minimum \hat{d}

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES

■ DPS

■ XES

New Algorithms

Results

Conclusions

Three source of heuristic information: h, \hat{h}, \hat{d}

EES search strategy:

$best_f$: open node giving lower bound on cost

$best_{\hat{f}}$: open node giving estimated optimal cost

$best_{\hat{d}}$: estimated ω -suboptimal node with minimum \hat{d}

node to expand next:

1. pursue the nearest goal estimated to lie within the bound
- 2.
- 3.

in other words:

1. **if** $\hat{f}(best_{\hat{d}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{d}}$
- 2.
- 3.

Three source of heuristic information: h, \hat{h}, \hat{d}

EES search strategy:

$best_f$: open node giving lower bound on cost

$best_{\hat{f}}$: open node giving estimated optimal cost

$best_{\hat{d}}$: estimated ω -suboptimal node with minimum \hat{d}

node to expand next:

1. pursue the nearest goal estimated to lie within the bound
2. pursue the estimated optimal solution
- 3.

in other words:

1. **if** $\hat{f}(best_{\hat{d}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{d}}$
2. **else if** $\hat{f}(best_{\hat{f}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{f}}$
- 3.

Three source of heuristic information: h, \hat{h}, \hat{d}

EES search strategy:

$best_f$: open node giving lower bound on cost

$best_{\hat{f}}$: open node giving estimated optimal cost

$best_{\hat{d}}$: estimated ω -suboptimal node with minimum \hat{d}

node to expand next:

1. pursue the nearest goal estimated to lie within the bound
2. pursue the estimated optimal solution
3. raise the lower bound on optimal solution cost

in other words:

1. **if** $\hat{f}(best_{\hat{d}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{d}}$
2. **else if** $\hat{f}(best_{\hat{f}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{f}}$
3. **else** $best_f$

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES

■ DPS

■ XES

New Algorithms

Results

Conclusions

Three source of heuristic information: h, \hat{h}, \hat{d}

EES search strategy:

1. **if** $\hat{f}(best_{\hat{d}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{d}}$
2. **else if** $\hat{f}(best_{\hat{f}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{f}}$
3. **else** $best_f$

Three source of heuristic information: h, \hat{h}, \hat{d}

EES search strategy:

1. **if** $\hat{f}(best_{\hat{d}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{d}}$
2. **else if** $\hat{f}(best_{\hat{f}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{f}}$
3. **else** $best_f$

Other EES variants:

1. **if** $\hat{f}(best_{\hat{d}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{d}}$
2. ~~**else if** $\hat{f}(best_{\hat{f}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{f}}$?~~
3. **else** $best_f$

see paper for more details.

Three source of heuristic information: h, \hat{h}, \hat{d}

EES search strategy:

1. **if** $\hat{f}(best_{\hat{d}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{d}}$
2. **else if** $\hat{f}(best_{\hat{f}}) < \omega \cdot f(best_f)$ **then** $best_{\hat{f}}$
3. **else** $best_f$

Problems:

- EES does not consider the uncertainty of its estimates (brittle)
- EES does not spend enough effort on estimating the bound

State-of-The-Art: 2/2 DPS (Gilon, Felner, and Stern, 2016)

Best first search on “potential”:

$$ud(n) = \frac{\omega \cdot f_{min} - g(n)}{h(n)}$$

[Introduction](#)

[Bounded Suboptimal](#)

■ Problem Setting

■ EES

■ DPS

■ XES

[New Algorithms](#)

[Results](#)

[Conclusions](#)

State-of-The-Art: 2/2 DPS (Gilon, Felner, and Stern, 2016)

Best first search on “potential”:

$$ud(n) = \frac{\omega \cdot f_{min} - g(n)}{h(n)}$$

does not explicitly optimize search time

[Introduction](#)

[Bounded Suboptimal](#)

■ Problem Setting

■ EES

■ DPS

■ XES

[New Algorithms](#)

[Results](#)

[Conclusions](#)

Bounded-Cost: XES (Fickert, Gu, and Ruml, 2021)

Best first search on **expected search effort**:

$$xe(n) = \frac{T(n)}{p(n)}$$

[Introduction](#)

[Bounded Suboptimal](#)

■ Problem Setting

■ EES

■ DPS

■ XES

[New Algorithms](#)

[Results](#)

[Conclusions](#)

Bounded-Cost: XES (Fickert, Gu, and Ruml, 2021)

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES

■ DPS

■ XES

New Algorithms

Results


Conclusions

Best first search on **expected search effort**:

$$xe(n) = \frac{T(n)}{p(n)}$$

$T(n)$: total search effort, estimated by $d(n)$

$p(n)$: the probability of finding a solution within the bound, estimated by:



Bounded-Cost: XES (Fickert, Gu, and Ruml, 2021)

Introduction

Bounded Suboptimal

■ Problem Setting

■ EES

■ DPS

■ XES

New Algorithms

Results


Conclusions

Best first search on **expected search effort**:

$$xe(n) = \frac{T(n)}{p(n)}$$

$T(n)$: total search effort, estimated by $d(n)$

$p(n)$: the probability of finding a solution within the bound, estimated by:



XES performs better than BEES and PS.

Can we adapt XES to bounded-suboptimal setting?

Introduction

Bounded Suboptimal

New Algorithms

■ DXES

■ RoundRobin

Results

Conclusions

New Algorithms

Our Approach: 1/2 Dynamic XES

Introduction

Bounded Suboptimal

New Algorithms

■ DXES

■ RoundRobin

Results


Conclusions

Best first search on **expected search effort**:

$$xe(n) = \frac{T_n}{p(n)}$$

$T(n)$: total search effort, estimated by $d(n)$

$p(n)$: the probability of finding a solution within the bound, estimated by:




hard to estimate when raising the bound is useful!

- Introduction
- Bounded Suboptimal
- New Algorithms
- DXES**
- RoundRobin
- Results
- Conclusions

$$xe(n) = \frac{T_n}{p(n)}$$

$p(n)$: the probability of finding a solution within the **estimated bound**, estimated by:



New Results in Bounded-Suboptimal Search – 13 / 21

Our Approach: 1/2 Dynamic XES

Introduction

Bounded Suboptimal

New Algorithms

■ DXES

■ RoundRobin

Results


Conclusions

Best first search on **expected search effort**:

$$xe(n) = \frac{T_n}{p(n)}$$

$T(n)$: total search effort, estimated by $d(n)$

$p(n)$: the probability of finding a solution within the **estimated bound**, estimated by:



hard to estimate when raising the bound is useful!

Our Approach: 1/2 Dynamic XES

Introduction

Bounded Suboptimal

New Algorithms

■ DXES

■ RoundRobin

Results


Conclusions

Best first search on **expected search effort**:

$$xe(n) = \frac{T_n}{p(n)}$$

$T(n)$: total search effort, estimated by $d(n)$

$p(n)$: the probability of finding a solution within the **estimated bound**, estimated by:



hard to estimate when raising the bound is useful!

Our Approach: 2/2 A Round-Robin Scheme

Introduction

Bounded Suboptimal

New Algorithms

■ DXES

■ RoundRobin

Results

Conclusions

Round-Robin on:

focal list: sorted by $d(\text{EES})$ or $ud(\text{DPS})$ or $xe(\text{DXES})$

open list: sorted by \hat{f}

cleanup list: sorted by f

focal and open condition: $f(n) < \omega \cdot f_{min}$

Simple but work well empirically!

Introduction

Bounded Suboptimal

New Algorithms

Results

■ Experiments

■ Planning

■ Search

Conclusions

Results

[Introduction](#)

[Bounded Suboptimal](#)

[New Algorithms](#)

[Results](#)

■ [Experiments](#)

■ [Planning](#)

■ [Search](#)

[Conclusions](#)

Planning Domains:

- Implementation in Fast Downward²
- Benchmarks:
IPS optimal tracks (48 domains)

Search Domains:

- Sliding-Tile Puzzle, Vacuum World, Pancake, Racetrack

²Helmert 2006.

IPC Coverage ($\omega = 1.5$)

- Introduction
- Bounded Suboptimal
- New Algorithms
- Results
 - Experiments
 - Planning
 - Search
- Conclusions

Coverage	WA [*]	EES	DPS	DXES	RR- <i>d</i>	RR-DPS	RR-DXES
Sum (1652)	995	967	1012	894	1025	982	1052
Normalized(%)	58.7	57.0	60.0	51.5	60.7	57.9	62.5
Expansions	569	558	472	734	383	665	371

→ RR-DXES and RR-*d* perform best overall.

Search Domains

[Introduction](#)

[Bounded Suboptimal](#)

[New Algorithms](#)

[Results](#)

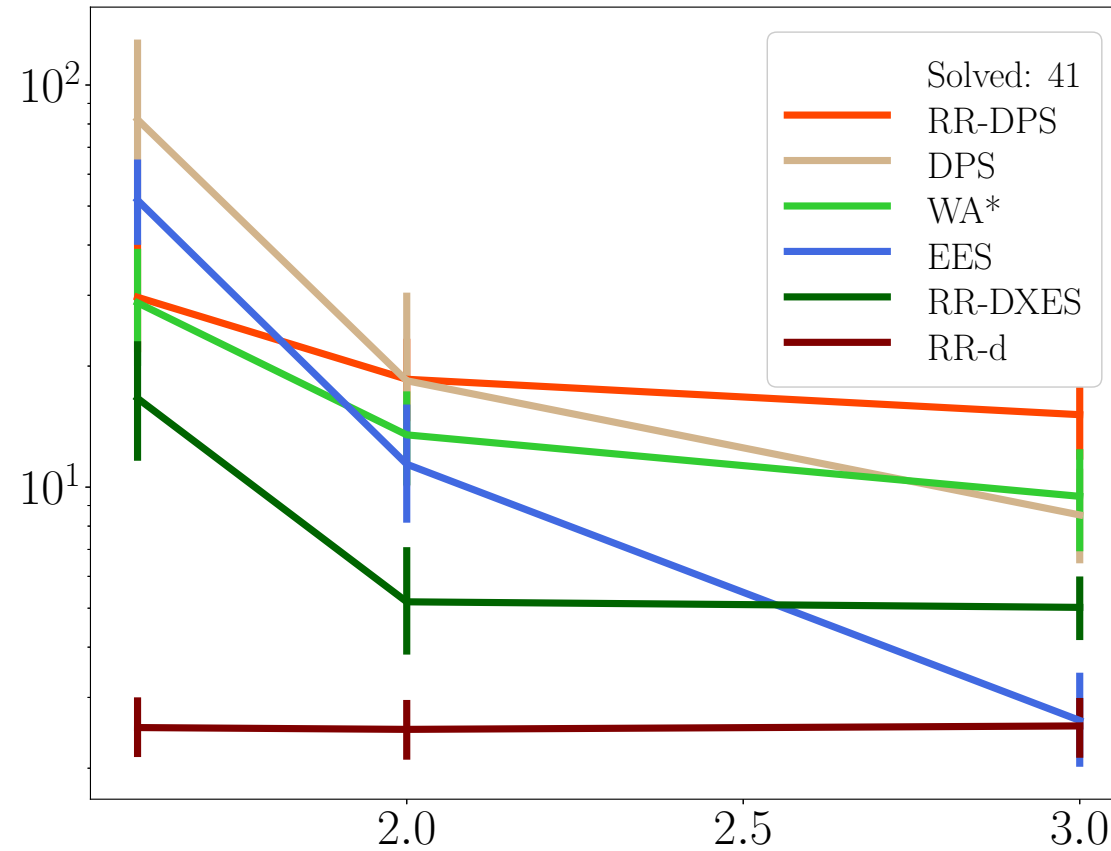
■ Experiments

■ Planning

■ Search

[Conclusions](#)

Uniform Vacuum World



Search Domains

[Introduction](#)

[Bounded Suboptimal](#)

[New Algorithms](#)

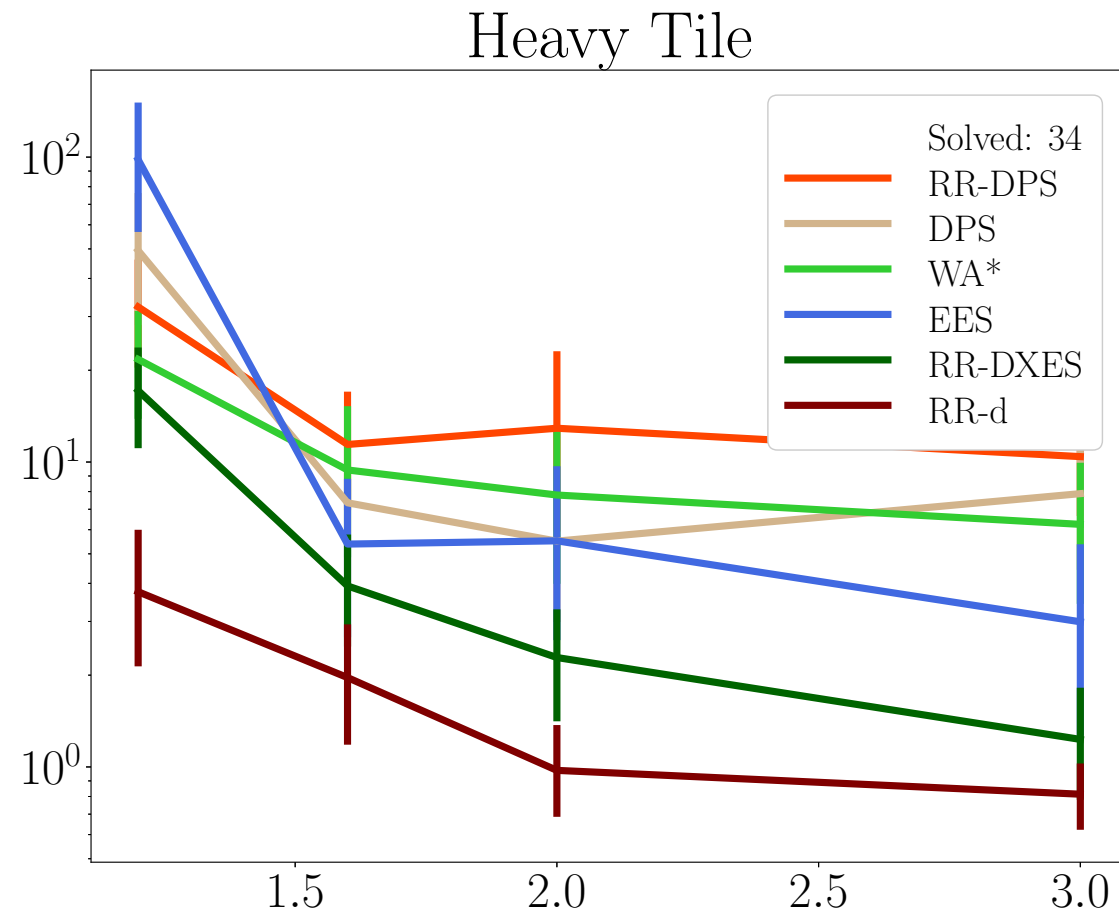
[Results](#)

■ Experiments

■ Planning

■ Search

[Conclusions](#)



Search Domains

[Introduction](#)

[Bounded Suboptimal](#)

[New Algorithms](#)


[Results](#)

■ Experiments

■ Planning

■ Search

[Conclusions](#)



[Introduction](#)

[Bounded Suboptimal](#)

[New Algorithms](#)

[Results](#)

[Conclusions](#)

■ Summary

Conclusions

Introduction

Bounded Suboptimal

New Algorithms

Results

Conclusions

■ Summary

What to do for bounded-suboptimal search:

- Weighted-A* is the first thing to try
- Round-Robin on d , \hat{f} , f is the next to try

What to do for bounded-suboptimal search:

- Weighted-A* is the first thing to try
- Round-Robin on d , \hat{f} , f is the next to try

What's next:

- When to raise bound, and when to pursue solution?
- DXES should help, but not quite, why?

Questions?

[Introduction](#)

[Bounded Suboptimal](#)

[New Algorithms](#)

[Results](#)

[Conclusions](#)

[Questions](#)

■ [Questions?](#)

